

# Threads

Juan José Costa and Alejandro Pajuelo


Escola Tècnica Superior de Telecomunicacions de Barcelona  
(ETSETB)

Universitat Politècnica de Catalunya (UPC)

2021-2022Q2

# Problem

- Cooperative processes send information to each other using system calls since they do not share memory

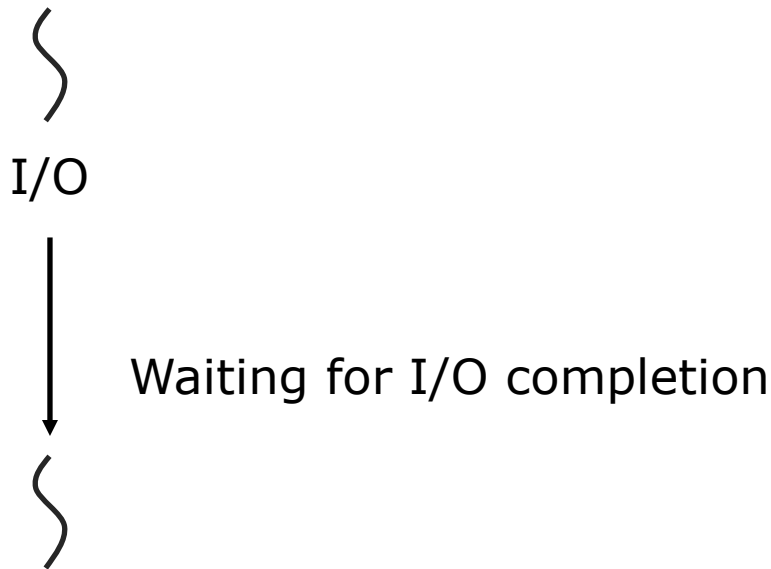


- Introduces overhead
- Parallel tasks can become sequential
- **Solution**: shared memory

# Problem

---

- I/O cannot be parallelized with just one flow of code execution



- Maybe, some computations could be done while waiting for I/O

# Threads - Why?

---

- When and why they are used...
  - Parallelism exploitation (code and hardware resources)
  - Encapsulate tasks (modular coding)
  - I/O efficiency (specific Threads for I/O)
  - Pipelining of service requests (to maintain QoS of services)
  
- Advantages
  - Creation/termination/context switches of threads are faster than with processes
  - With memory sharing inside the process, threads can share information without system calls

# Why not?



**Davidlohr Bueso**

@davidlohr



A programmer had a problem. He thought to himself, "I know, I'll solve it with threads!". has Now problems. two he

12:16 AM · Jan 9, 2013



1.5K



4.5K




Copy link to Tweet

## ■ Disadvantages

- Difficult coding and debugging due to the shared memory
  - ▶ Problems of synchronization and mutual exclusion
    - Incoherent executions, wrong results, deadlocks, ...

# Threads

- A process has one thread when execution starts
- A process can have several threads
  - High performance videogames **>50 threads**;
  - Firefox/Chrome **>80 threads**
- In the next Figure: Process 1 has 2 threads; Process 2 has 4 threads; Process 3 has 1 thread
- Process management with several threads depends on the OS support
  - **User Level Threads** vs **Kernel Level Threads**



# What are threads?


---

- Concept of a process:
  - OS representation of a program in execution
  - Resource container
  
- A thread is one of the resources of a process
  - Minimal scheduling unit
    - ▶ Every independent part of the code can be assigned to a thread
  - Thread resources
    - ▶ Identifier (Thread ID: TID)
    - ▶ Stack Pointer
    - ▶ Program Counter
    - ▶ Register File
    - ▶ Errno
  - Threads share the resources of the process they belong to

# Thread sample

---

```
int global = 2;
int f1(void) {
    ...
}
int f2(void) {
    ...
}
int main() {
    f1()
    f2()
    do_things();
}
```

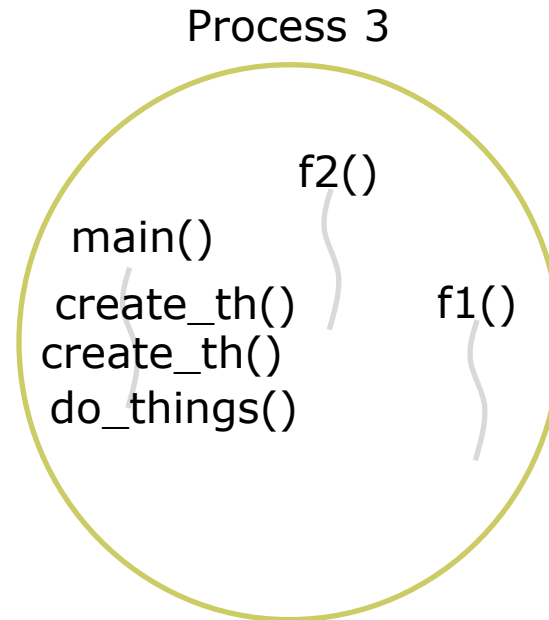




# Thread sample

---

```
int global = 2;
int f1(void) {
    ...
}
int f2(void) {
    ...
}
int main() {
    create_tread( f1 )
    create_thread( f2 )
    do_things();
}
```



# Basic services (UNIX)



Servicio	Llamada a sistema
Thread creation	pthread_create
Thread finalization	pthread_exit
Wait for a thread	pthread_join
Identification	pthread_self
Scheduling	pthread_yield

- ❑ Do not forget to include the header:  
`#include <pthread.h>`
- ❑ And link with `-pthread`

# Thread creation

---

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine) (void *), void *arg);
```

- `pthread_t *thread`: identifier of the created thread
- `pthread_attr_t *attr`: initial attributes to create the thread
- `void *(*start_routine)(void *)`: function with the thread code
- `void *arg`: argument to pass to thread

```
start_routine(arg);
```

- Thread identifiers are only valid inside the process (`pthread_self`)
- Don't forget to link with `-lpthread`

# Thread creation: example

---

```
void * func(void *arg)
{
}

int main()
{
    ...
    pthread_t id;
    pthread_create(&id, NULL, func, NULL);
    ...
}
```

# Thread creation: example

---

## ■ Creation of 10 threads

```
void * func(void *arg)
{
}

int main()
{
    ...
    int i;
    pthread_t id[10];
    for (i=0; i<10; i++)
    {
        pthread_create(&id[i], NULL, func, NULL);
    }
    ...
}
```

# Thread creation: example

---

## ■ Argument passing: creation order

```
void * func(void *arg)
{
    int order = (int)arg;
}

int main()
{
    ...
    int i;
    pthread_t id[10];
    for (i=0; i<10; i++)
    {
        pthread_create(&id[i], NULL, func, (void*)i);
    }
    ...
}
```

# Thread termination

---

```
void pthread_exit(void *retval);
```

- void \*retval: termination code
- The compiler **implicitly** adds a call to pthread\_exit in case the programmer does not put it in the code

# Thread termination: example

---

```
void * func(void *arg)
{
    pthread_exit(0);
}

int main()
{
    ...
    pthread_t id;
    pthread_create(&id, NULL, func, NULL);
    ...
}
```



# Wait for a thread

---

```
int pthread_join(pthread_t thread, void **retval);
```

- `pthread_t thread`: identifier of the thread to wait for
- `void **retval`: termination code of the thread
- **Blocking** call
- `pthread_t thread` is the identifier returned when the thread was created
- `void **retval` is the value used in `pthread_exit`

# Wait for a thread: example

---

```
void * func(void *arg)
{
    pthread_exit(0);
}

int main()
{
    ...
    pthread_t id;
    int a;
    pthread_create(&id, NULL, func, NULL);
    pthread_join(id, (void**)&a);
    ...
}
```

# Wait for 10 threads (parallel): example

---

```
void * func(void *arg){
    int order = (int)arg;
    pthread_exit((void*)order);
}

int main(){
    ...
    int i;
    pthread_t id[10];
    int ret_val;
    for (i=0; i<10; i++){
        pthread_create(&id[i], NULL, func, (void*)i);
    }
    for (i=0; i<10; i++){
        pthread_join(id[i], (void**)&ret_val);
    }
    ...
}
```

# Wait for 10 threads (sequential): example

---

```
void * func(void *arg) {
    int order = (int)arg;
    pthread_exit((void*)order);
}

int main() {
    ...
    int i;
    pthread_t id[10];
    int ret_val;
    for (i=0; i<10; i++){
        pthread_create(&id[i], NULL, func, (void*)i);
        pthread_join(&id[i], (void**)&ret_val);
    }
    ...
}
```

# Specialization: example

---

- Create 1 thread that executes a CPU-intensive workload and overlap this computation with I/O

```
void * CPU_thread(void *arg)
{
    int order = (int)arg;
    Do_a_nice_computation_here();
    pthread_exit((void*)order);
}

void * IO_thread(void *arg)
{
    int order = (int)arg;
    Do_a_nice_IO_here();
    pthread_exit((void*)order);
}

int main()
{
    ...
    pthread_t id[2];
    int ret_val;
    pthread_create(&id[0], NULL, CPU_thread, (void*)1);
    pthread_create(&id[1], NULL, IO_thread, (void*)2);
    pthread_join(&id[0], (void**)&ret_val);
    pthread_join(&id[1], (void**)&ret_val);
    ...
}
```

# Shared Memory

---

- Threads:
  - Share the memory of the process they belong to
  - In general, they share all the resources assigned to the process
  
- Processes:
  - Don't share memory
  - Every process has its own memory

# Shared memory with threads: example

---

```
int a=0;


void * Thread_1(void *arg)
{
    while (1) a=1;
}

void * Thread_2(void *arg)
{
    while (1) a=2;
}

int main()
{
    ...
    pthread_t id[2];
    int ret_val;
    pthread_create(&id[0], NULL, Thread_1, NULL);
    pthread_create(&id[1], NULL, Thread_2, NULL);
    pthread_join(&id[0], (void**)&ret_val);
    pthread_join(&id[1], (void**)&ret_val);
    ...
}
```

# Shared memory with threads: example

---

**global variable** 

```
int a=0;

void * Thread_1(void *arg)
{
    while (1) a=1;
}

void * Thread_2(void *arg)
{
    while (1) a=2;
}

int main()
{
    ...
    pthread_t id[2];
    int ret_val;
    pthread_create(&id[0], NULL, Thread_1, NULL);
    pthread_create(&id[1], NULL, Thread_2, NULL);
    pthread_join(&id[0], (void**)&ret_val);
    pthread_join(&id[1], (void**)&ret_val);
    ...
}
```



# Shared memory with threads: example

```
int a=0;
void * Thread_1(void *arg)
{
    while (1) a=1;
}
void * Thread_2(void *arg)
{
    while (1) a=2;
}
int main()
{
    ...
    pthread_t id[2];
    int ret_val;
    pthread_create(&id[0], NULL, Thread_1, NULL);
    pthread_create(&id[1], NULL, Thread_2, NULL);
    pthread_join(&id[0], (void**)&ret_val);
    pthread_join(&id[1], (void**)&ret_val);
    ...
}
```

global variable

(visible by all threads)

# Shared memory with threads: example

```
int a=0;
void * Thread_1(void *arg)
{
    while (1) a=1;
}

void * Thread_2(void *arg)
{
    while (1) a=2;
}

int main()
{
    ...
    pthread_t id[2];
    int ret_val;
    pthread_create(&id[0], NULL, Thread_1, NULL);
    pthread_create(&id[1], NULL, Thread_2, NULL);
    pthread_join(&id[0], (void**)&ret_val);
    pthread_join(&id[1], (void**)&ret_val);
    ...
}
```

global variable

(visible by all threads)

What happens if I put  
a loop here printing  
var 'a' value?

# Shared memory with processes: example

---

## ■ Processes do not share memory!!!!

```
int a=0;

void * Proc1_1()
{
    while (1) a=1;
}

void * Proc_2()
{
    while (1) a=2;
}

int main()
{
    ...
    if (fork()==0) Proc_1();
    if (fork()==0) Proc_2();
    ...
}
```

# Shared memory with processes: example

---

## ■ Processes do not share memory!!!!

global variable

```
int a=0;
void * Proc1_1()
{
    while (1) a=1;
}

void * Proc_2()
{
    while (1) a=2;
}

int main()
{
    ...
    if (fork()==0) Proc_1();
    if (fork()==0) Proc_2();
    ...
}
```

# Shared memory with processes: example

## ■ Processes do not share memory!!!!

**global variable**

```
int a=0;
void * Proc1_1()
{
    while (1) a=1;
}

void * Proc_2()
{
    while (1) a=2;
}

int main()
{
    ...
    if (fork()==0) Proc_1();
    if (fork()==0) Proc_2();
    ...
}
```

**(visible by all processes)**

# Shared memory with processes: example

## ■ Processes do not share memory!!!!

**global variable**

```
int a=0;
void * Proc1_1()
{
    while (1) a=1;
}

void * Proc_2()
{
    while (1) a=2;
}

int main()
{
    ...
    if (fork()==0) Proc_1();
    if (fork()==0) Proc_2();
    ...
}
```

**(visible by all processes)**

**What happens if I put a loop here printing var 'a' value?**

# Scheduler

---

- If an OS supports threads, the scheduler must be modified:
  - Processes are not longer scheduled since they are resource containers
- Two ways:
  - 1 level scheduling: thread scheduling
    - ▶ All the threads in the system are considered
  - 2 level scheduling:
    - ▶ The next thread to execute is chosen among the ready threads of the current process
    - ▶ In case that no thread is candidate, select one thread from another process
- Algorithms and scheduling policies are the same

# Relinquish the CPU

---

- There are times where the CPU is not needed and you may voluntarily free the CPU for others threads to execute.

```
int pthread_yield( void );
```

- Relinquishes the CPU and another thread is scheduled to run.
- (For processes the sched\_yield is also available)



# Summary

---

- Thread concepts
- Services
  - Creation → `pthread_create`
  - Finalization → `pthread_exit`, `pthread_join`
  - Identification → `pthread_self`
  - Scheduling → `pthread_yield`