

Communication and Synchronization between tasks

Juan José Costa and Alejandro Pajuelo

Escola Tècnica Superior de Telecomunicacions de Barcelona
(ETSETB)

Universitat Politècnica de Catalunya (UPC)

2021-2022Q2

Index

- Why communicate and synchronize?
- Shared memory
- Distributed memory

Introduction

- A single process is easy, but no practical
- More than one process is desirable
 - To cooperate in some problem
 - To accelerate some computation
 - For modularity (different roles)
- But then new needs arise:
 - How do we pass info from one to the other?
 - What if we want to execute one before the other?

Introduction

- We need ways to:
 - share information
 - synchronize processes
- We will see different ways to communicate processes
 - Inter Process Communication (IPC)
- There are 2 communication models:
 - Shared memory
 - Processes use some variables that can be read/written
 - Message passing
 - Processes use specific functions to send/receive data

Index

- Why communicate and synchronize?
- Shared memory
- Distributed memory

Shared Memory

- Introduction
- Problems
 - Critical Region
 - Serialization
- Mutual exclusion
 - Busy-waiting
 - Token [Harder18]sec 9.3 [sync]
 - Test-Set [Harder18]sec 9.4
 - Mutex [Harder18]sec 9.5.1
 - Semaphores [Harder18]sec 9.5 [sync]

Introduction

- A single machine has the following resources:
 - memory
 - secondary memory
 - network
- They may be shared by different tasks

Introduction: Example

```
int first = 1; /* GLOBAL VAR */
```

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```


Introduction: Example

```
int first = 1; /* GLOBAL VAR */
```

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

With a single task we are fine 😊

- No sharing at all

Introduction: Example

```
int first = 1; /* GLOBAL VAR == SHARED */  
create_2_threads( f );
```

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

Introduction: Example

```
int first = 1; /* GLOBAL VAR == SHARED */  
create_2_threads( f );
```

1st thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

2nd thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

Possible outcomes?

Introduction: Example

```
int first = 1; /* GLOBAL VAR == SHARED */  
create_2_threads( f );
```

1st thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

2nd thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

Possible outcomes?

task1()	task2()
1 st thread	2 nd thread

Introduction: Example

```
int first = 1; /* GLOBAL VAR == SHARED */  
create_2_threads( f );
```

1st thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

2nd thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

Possible outcomes?

task1()	task2()
1 st thread	2 nd thread
2 nd thread	1 st thread

Introduction: Example

```
int first = 1; /* GLOBAL VAR == SHARED */  
create_2_threads( f );
```

1st thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

2nd thread

```
int f() {  
    if (first) {  
        first --;  
        task1();  
    }  
    else {  
        task2();  
    }  
}
```

Possible outcomes?

task1()	task2()
1 st thread	2 nd thread
2 nd thread	1 st thread
1 st and 2 nd	

Problem

- Shared variable
 - `first`
- Non atomic operations:
 - LOAD
 - SUB
 - STORE
- **Concurrent accesses to shared resources** can lead to unexpected or erroneous behavior
 - Known as *race conditions*
- Another example: [Harder18] Section 9.1.1

Critical Region

- **Section of a program** where a shared resource is accessed concurrently and needs to be protected
- This section is the **critical region**.
- A mechanism is needed to ensure that it can only be executed by a single task at a time.
 - **Mutual exclusion**

Two tasks communicating information

- Two iterative tasks trying to share information.
- Task 1 (*Producer*) prepares data
 - which is to be sent to and used by Task 2.
- Data should not be overwritten by Task 1
 - until Task 2 has completed using it.
- Task 2 (*Consumer*) should not access data
 - until Task 1 has finished writing to it.
- (Courtesy of Harder18 Sec. 9.1)

Two tasks communicating information

```
Data *result;

// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result


        write( result );

        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...

        read( result );

        // Continue executing...
    }
}
```



Two tasks communicating information:

1st try

```
#include <stdbool.h>
```

```
bool result_is_produced = false;  Add a flag
```

```
Data *result;
```

```
// Producer
```

```
void task_1( void ) {
```

```
    while ( 1 ) {
```

```
        // Executing...
```

```
        // Prepare something for
```

```
        // task 2 and assign to result
```

```
        write( result );
```

```
        result_is_produced = true;
```

```
        // Continue executing...
```

```
    }
```

```
}
```

```
// Consumer
```

```
void task_2( void ) {
```

```
    while( 1 ) {
```

```
        // Executing...
```

```
        while( !result_is_produced );
```

```
        read( result );
```

```
        // Continue executing...
```

```
    }
```

```
}
```

Two tasks communicating information:

1st try

```
#include <stdbool.h>
```

```
bool result_is_produced = false;  Add a flag
```

```
Data *result;
```

```
// Producer
```

```
void task_1( void ) {
```

```
    while ( 1 ) {
```

```
        // Executing...
```

```
        // Prepare something for
```

```
        // task 2 and assign to result
```

```
        while( result_is_produced );
```

```
        write( result );
```

```
        result_is_produced = true;
```

```
        // Continue executing...
```

```
    }
```

```
}
```

```
// Consumer
```

```
void task_2( void ) {
```

```
    while( 1 ) {
```

```
        // Executing...
```

```
        while( !result_is_produced );
```

```
        read( result );
```

```
        result_is_produced = false;
```

```
        // Continue executing...
```

```
    }
```

```
}
```

Two tasks communicating information:

1st try

Different case as before
(both codes are mutually exclusive)

```
#include <stdbool.h>
bool result_is_produced = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced );
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}
```

```
// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced );
        read( result );
        result_is_produced = false;
        // Continue executing...
    }
}
```

Add a flag

Two tasks communicating information:

1st try

```
#include <stdbool.h>
bool result_is_produced = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced );
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced );
        read( result );
        result_is_produced = false;
        // Continue executing...
    }
}
```

What happens in a single-processor machine?

Two tasks communicating information:

1st try

```
#include <stdbool.h>
bool result_is_produced = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced );
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced );
        read( result );
        result_is_produced = false;
        // Continue executing...
    }
}
```

What happens in a single-processor machine? **Busy Waiting**

Two tasks communicating information:

2nd try

```
#include <stdbool.h>
bool result_is_produced = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced ) {
            pthread_yield();
        }
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced ) {
            pthread_yield();
        }
        read( result );
        result_is_produced = false;
        // Continue executing...
    }
}
```


Two tasks communicating information:

2nd try

```
#include <stdbool.h>
bool result_is_produced = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced ) {
            pthread_yield();
        }
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced ) {
            pthread_yield();
        }
        read( result );
        result_is_produced = false;
        // Continue executing...
    }
}
```

Round Robin... ok, but, priorities?...

Multiple tasks communicating information

- Imagine 2 copies of Task2
 - they can NOT access data structure at same time

Multiple tasks communicating information

```
#include <stdbool.h>
bool result_is_produced = false;
bool consumer_is_reading = false;
Data *result;

// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced ) {
            pthread_yield();
        }
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}

// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced
            || consumer_is_reading ) {
            pthread_yield();
        }
        consumer_is_reading = true;
        read( result );
        result_is_produced = false;
        consumer_is_reading = false;
        // Continue executing...
    }
}
```

Multiple tasks communicating information

```
#include <stdbool.h>
bool result_is_produced = false;
bool consumer_is_reading = false;
Data *result;
// Producer
void task_1( void ) {
    while ( 1 ) {
        // Executing...
        // Prepare something for
        // task 2 and assign to result
        while( result_is_produced ) {
            pthread_yield();
        }
        write( result );
        result_is_produced = true;
        // Continue executing...
    }
}
```

Sharing problem! ☹️

```
// Consumer
void task_2( void ) {
    while( 1 ) {
        // Executing...
        while( !result_is_produced
            || consumer_is_reading ) {
            pthread_yield();
        }
        consumer_is_reading = true;
        read( result );
        result_is_produced = false;
        consumer_is_reading = false;
        // Continue executing...
    }
}
```

Summary

- We have seen 2 different problems:
 1. **Mutual exclusion:** preventing more than one task from accessing data, and
 2. **Serialization:** ensuring that one task is performed after another.

Shared Memory

- Introduction
- Problems
 - Critical Region
 - Serialization
- Mutual exclusion
 - Busy-waiting
 - Token [Harder18]sec 9.3 [sync]
 - Test-Set [Harder18]sec 9.4
 - Mutex [Harder18]sec 9.5.1
 - Semaphores [Harder18]sec 9.5 [sync]

Mutual Exclusion

- Protection mechanism for a critical region
- Ensure that no other thread may execute the region while a thread is inside
 - Even in presence of context switches!

Mutual Exclusion: Token passing

- A simple way: force thread execution order
- Token passing mechanism
 - Identify each consumer
 - Generate a token with one of the identifiers
 - When a consumer has accessed the result
 - update token to the identifier of the next consumer

Token passing example

- Define token structure:

```
typedef struct {  
    size_t this_id;  
    size_t next_id;  
} token_t;
```

- Establish the order (3 tasks):

```
token_t args[3] = {{1, 2}, {2, 3}, {3, 1}};  
// 1->2->3->1->2->...
```

- Initialize the starting identifier task and tasks:

```
size_t reading = 1;
```

```
pthread_create( &thread_id[0], NULL, &consumer, &args[0] );  
pthread_create( &thread_id[1], NULL, &consumer, &args[1] );  
pthread_create( &thread_id[2], NULL, &consumer, &args[2] );
```

Token passing example

- Then the consumer becomes:

```
void *consumer( void *p_arguments ) {
    token_t *p_identifier = (token_t *)p_arguments;
    while ( 1 ) {
        while ( ( reading != p_identifier->this_id ) ||
                !result_is_produced ) {
            pthread_yield();
        }
        --result;
        printf( "%d ", result );
        result_is_produced = false;
        reading = p_identifier->next_id;
    }
}
```

Mutual Exclusion

- Allow any order
- Programmer needs to mark the start and end of the critical region:
 - **start**: if there is no thread in the critical region, enter; otherwise wait until the working thread finishes
 - **end**: release access to the critical region; if any other thread was waiting, allow it to enter

Mutual exclusion: 1st try

```
bool ready = true; //Global variable

// Start Critical Region
while ( !ready ) {
    scheduler();
}
ready = false;
// Inside the critical region!
// Access the data structure(Mutual exclusion)
// End critical region
ready = true;
// Out of Critical Region
```

Mutual exclusion: 1st try

```
bool ready = true; //Global variable
```

```
// Start Critical Region  
while ( !ready ) {  
    scheduler();  
}  
ready = false;
```


```
// Inside the critical region!  
// Access the data structure(Mutual exclusion)
```

```
// End critical region  
ready = true;  
// Out of Critical Region
```

Mutual exclusion: 1st try

```
bool ready = true; //Global variable
```

```
// Start Critical Region  
while ( !ready ) {  
    scheduler();  
}
```



```
ready = false;
```

```
// Inside the critical region!  
// Access the data structure(Mutual exclusion)
```

```
// End critical region  
ready = true;
```

```
// Out of Critical Region
```

- Same problem as before ☹
- We must have some means of test and set a variable
 - so that no interrupt can occur between both

Test-and-Reset*

- We want a function that if parameter is
 - *false*, it returns *false*,
 - *true*, it is set to *false*, but returns the old value (*true*).
- The function `test_and_reset(...)` is a function that is translated to a single machine instruction which can be thought of as:

```
bool test_and_reset( bool *value ) {  
    bool previous = *value;  
    *value = false;  
    return previous;  
}
```

Mutual exclusion: 2nd try

```
bool ready = true; //Global variable

// Enter Critical Region
while ( !test_and_reset(&ready) ) {
    scheduler();
}

// Inside the critical region!
// Access the data structure(Mutual exclusion)
// Exit critical region
ready = true;
// End Critical Region
```


Mutual exclusion: 2nd try

```
bool ready = true; //Global variable
```

mutex_init

```
// Enter Critical Region
while ( !test_and_reset(&ready) ) {
    scheduler();
}

// Inside the critical region!
// Access the data structure(Mutual exclusion)
// Exit critical region
ready = true;
// End Critical Region
```

Mutual exclusion: 2nd try

```
bool ready = true; //Global variable
```

mutex_init

```
// Enter Critical Region  
while ( !test_and_reset(&ready) ) {  
    scheduler();  
}
```

mutex_lock

```
// Inside the critical region!  
// Access the data structure(Mutual exclusion)  
// Exit critical region  
ready = true;  
// End Critical Region
```

Mutual exclusion: 2nd try

```
bool ready = true; //Global variable
```

mutex_init

```
// Enter Critical Region  
while ( !test_and_reset(&ready) ) {  
    scheduler();  
}
```

mutex_lock

```
// Inside the critical region!
```

```
// Access the data structure(Mutual exclusion)
```

```
// Exit critical region
```

```
ready = true;
```

mutex_unlock

```
// End Critical Region
```

Mutual exclusion: 2nd try

```
bool ready = true; //Global variable
```

mutex_init

```
// Enter Critical Region  
while ( !test_and_reset(&ready) ) {  
    scheduler();  
}
```

mutex_lock

```
// Inside the critical region!
```

```
// Access the data structure(Mutual exclusion)
```

```
// Exit critical region
```

```
ready = true;
```

mutex_unlock

```
// End Critical Region
```

- But real-time systems are tricky and if one of the tasks is high priority ...

Mutual Exclusion: Example

```
int first = 1; /* GLOBAL VAR == SHARED */
mutex_t critical;
mutex_init(&critical)
create_2_threads( f );
```

1st thread

```
int f() {
    mutex_lock(&critical)
    if (first) {
        first --;
        mutex_unlock(&critical)
        task1();
    } else {
        mutex_unlock(&critical)
        task2();
    }
}
```

2nd thread

```
int f() {
    mutex_lock(&critical)
    if (first) {
        first --;
        mutex_unlock(&critical)
        task1();
    } else {
        mutex_unlock(&critical)
        task2();
    }
}
```

Pthread Mutex API

- **#include <pthread.h>**
- `pthread_mutex_t fastmutex = PTHREAD_MUTEX_INITIALIZER;`
- `int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *mutexattr);`
- `int pthread_mutex_lock(pthread_mutex_t *mutex);`
- `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
- `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
- `int pthread_mutex_destroy(pthread_mutex_t *mutex);`

Pthread Mutex example

- A couple of threads increasing a global var

Pthread Mutex example

- A couple of threads increasing a global var

```
pthread_mutex_t critical; // Declare a mutex global
```


Pthread Mutex example

- A couple of threads increasing a global var

```
pthread_mutex_t critical; // Declare a mutex global

main() {
    pthread_mutex_init(&critical, NULL); // BEFORE thread creation
    ...
    pthread_create( ... );
    pthread_create( ... );
}
```

Pthread Mutex example

- A couple of threads increasing a global var


```
pthread_mutex_t critical; // Declare a mutex global

main() {
    pthread_mutex_init(&critical, NULL); // BEFORE thread creation
    ...
    pthread_create( ... );
    pthread_create( ... );
}

pthread_mutex_lock(&critical); // Enter critical region
pthread_mutex_unlock(&critical); // Exit critical region
```

Dining philosophers

- think → pick left fork → pick right fork → eat → drop right fork → drop left fork → think → ...
- Deadlock



https://en.wikipedia.org/wiki/Dining_philosophers_problem

Deadlock

- The following rules must comply to generate a deadlock:
 - Mutual exclusion: At least 2 non shareable resources
 - Hold & Wait: You get a resource and wait for another
 - No preemption: No one can remove my resource (just me)
 - Circular wait: There is circular sequence of 2 or more tasks, where each one needs a resource being used by another one.

Breaking deadlocks

- Break ANY of the previous rules:
 - Add shareable resources
 - Enable resource stealing
 - Pick all resources or none
 - Sort the requests (enforce the order in which the resources are requested)

POSIX semaphores API

- POSIX semaphores allow processes and threads to **synchronize** their actions.
 - Unnamed → shared memory
 - Named → shared filesystem
- A semaphore is an integer whose value is never allowed to fall below zero.
- Two operations can be performed on semaphores:
 - **decrement** the semaphore value by one (**sem_wait**).
 - If value is 0 → will block until the value becomes greater than zero.
 - and **increment** the semaphore value by one (**sem_post**)
 - if value becomes >0 → unblock a previously blocked proc/thread
- (Courtesy Linux Programmer's Manual)

POSIX semaphores(unnamed)

- **#include <semaphore.h>**
- **int sem_init(sem_t *sem, int pshared, unsigned int value);**
- **int sem_wait(sem_t *sem);**
- **int sem_post(sem_t *sem);**
- **int sem_trywait(sem_t *sem);**
- **int sem_timedwait(sem_t *sem, const struct timespec *abs_timeout);**
- **int sem_destroy(sem_t *sem);**

Semaphores use

- `int sem_init(sem_t *sem, int pshared, unsigned int value);`
- The **value** used to initialize the semaphore defines its use:
 - **1**: MUTEX. Allow a single thread
 - **0**: SYNCHRONIZATION
 - **N**: RESTRICT RESOURCE USAGE. Allow only N threads
- **pshared** defines the sharing of the semaphore
 - 0: Only shared between the threads of a process
 - 1: Shared between processes

POSIX semaphore example

- A couple of threads increasing a global var

POSIX semaphore example

- A couple of threads increasing a global var

```
sem_t critical; // Declare a semaphore global
```

POSIX semaphore example

- A couple of threads increasing a global var

```
sem_t critical; // Declare a semaphore global

main() {
    sem_init(&critical, 0, 1); // BEFORE thread creation
                               // A MUTEX

    ...
    pthread_create( ... );
    pthread_create( ... );
}
```

POSIX semaphore example

- A couple of threads increasing a global var

```
sem_t critical; // Declare a semaphore global

main() {
    sem_init(&critical, 0, 1); // BEFORE thread creation
                               // A MUTEX

    ...
    pthread_create( ... );
    pthread_create( ... );
}

sem_wait(&critical); // Enter critical region
sem_post(&critical); // Exit critical region
```

POSIX semaphore example

- Execute in order.

POSIX semaphore example

- Execute in order.

```
sem_t s1; // Declare a semaphore per thread  
sem_t s2;
```

POSIX semaphore example

- Execute in order.

```
sem_t s1; // Declare a semaphore per thread
sem_t s2;

main() {
    sem_init(&s1, 0, 0); // BEFORE thread creation
                        // A SYNCHRONIZATION

    ...
    pthread_create( ... );
    pthread_create( ... );
    sem_post(&s1); // Wake 1st
}
```

POSIX semaphore example

- Execute in order.

```
sem_t s1; // Declare a semaphore per thread
sem_t s2;


main() {
    sem_init(&s1, 0, 0); // BEFORE thread creation
                        // A SYNCHRONIZATION

    ...
    pthread_create( ... );
    pthread_create( ... );
    sem_post(&s1); // Wake 1st
}

sem_wait(&s1); // Block at the beginning of the thread
sem_post(&s2); // Unblock following thread
```




Other ways to share memory

- Global variables are simple ways to share memory by related threads, but what for processes?




Other ways to share memory

- Global variables are simple ways to share memory by related threads, but what for processes?



Other ways to share memory

- Global variables are simple ways to share memory by related threads, but what for processes?



POSIX semaphores (named)

- `#include <semaphore.h>`
- `sem_t *sem_open(const char *name, int oflag);`
- `sem_t *sem_open(const char *name, int oflag,
mode_t mode, unsigned int value);`
- `int sem_close(sem_t *sem);`
- `int sem_unlink(const char *name);`

POSIX named semaphore example

1. Create the semaphore

```
sem_t *s = sem_open("mysemaphore", O_CREAT|O_RDWR, 0777, 0);
```

2. Open the semaphore (from other process)

```
sem_t *s = sem_open("mysemaphore", O_RDWR);
```

3. Use normally

4. Destroy after use

```
sem_close(s);  
sem_unlink("mysemaphore");
```

POSIX shared memory

```
#include <sys/mman.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <sys/types.h>
#include <unistd.h>
```

- `int shm_open (const char *name, int oflag, mode_t mode);`
- `int ftruncate (int fd, off_t length);`
- `void *mmap (void *addr, size_t length, int prot, int flags, int fd, off_t offset);`
- `int munmap (void *addr, size_t length);`
- `int shm_unlink (const char *name);`
- `int close (int fd);`
- Link with **-lrt**

POSIX shared memory example

1. Create the shared memory region

```
int r = shm_open("regionname", O_CREAT|O_RDWR, 0777, 0);  
ftruncate(r, 4096); //Set region's size
```

2. Map the region into process

```
int *region = (int*) mmap(NULL, 4096, PROT_READ|PROT_WRITE,  
MAP_SHARED, r, 0);  
close(r); // file descriptor may be closed now
```

3. Use normally

```
region[0] = xxxx  
region[1] = ...
```

4. Destroy after use

```
munmap(region, 4096)  
shm_unlink("regionname");
```

POSIX shared memory example

1. Open the shared memory region

```
int r = shm_open("regionname", O_RDWR, 0777, 0);
```

2. Map the region into process

```
int *region = (int*) mmap(NULL, 4096, PROT_READ|PROT_WRITE,  
MAP_SHARED, r, 0);  
close(r); // file descriptor may be closed now
```


3. Use normally

```
region[0] = xxxx  
region[1] = ...
```


4. Unmap after use

```
munmap(region, 4096);
```



POSIX shared memory example




POSIX shared memory example




POSIX shared memory example



POSIX shared memory example



POSIX shared memory example



Exercises

- Detect critical regions
- Protect region with mutex
- Protect region with semaphores
- Petr Hudeček and Michal Pokorný. The Deadlock Empire. HackCambridge 2016.
<https://deadlockempire.github.io/>

Shared memory Summary

- Problems:
 - Critical region
 - Serialization
- Solutions
 - Token passing
 - Mutual exclusion
 - Semaphores
- Sharing memory between processes

Index

- Why communicate and synchronize?
- Shared memory
- Distributed memory

Distributed memory

- 2 process may not share memory
 - same or different machines
 - [Harder18]ch12
- signals
- pipes
- message queues (Bertolotti12) 7
- sockets
- mailboxes (Bertolotti12)7
- Blocking operations

Considerations

- Message type
 - synchronization
 - communication
- Information of message
 - implicit
 - explicit
- Example: Post-office service 12.1

Considerations


- 2 main operations
 - send
 - receive
- Direct or indirect naming (Bertolotti12 ch 6)
 - mailboxes
- Synchronous or asynchronous

Distributed memory index

- Signals
- File Descriptors
- Pipes
- Message queues
- Sockets
- Mailboxes

Signals

- Signal: A notification stating that an event has happened
- May be sent by the kernel or same user tasks



Signals

- Each signal has an associated **handler**
 - That manages the reaction to it (behavior)
 - Predefined by the kernel
- Each signal has an identification number
- A task may capture/change their behavior
 - There are exceptions: SIGKILL and SIGSTOP
- Each task may have different handlers

Signals types

- Some signals:

Name	Default action	Event
SIGCHLD	IGNORE	A child has terminated or stopped
SIGCONT	CONTINUE	Continue if stopped
SIGSTOP	STOP	Stop task
SIGINT	TERMINATE	Keyboard interrupt (Ctrl-C)
SIGALRM	TERMINATE	The counter defined in 'alarm' call has finished
SIGKILL	TERMINATE	Finish the task
SIGSEGV	CORE	Invalid reference to memory
SIGUSR1	TERMINATE	User defined
SIGUSR2	TERMINATE	User defined

Signal management

- Works like a hardware interrupt generated by sw
 - When a task receives a signal
 - interrupt current code execution
 - jump to signal management code (handler)
 - if the handler does not kill the process
 - return to previous code execution
- A task may block/unblock the reception of signals
 - It has a **mask of blocked signals**
 - Except SIGKILL and SIGSTOP
 - On reception of a blocked signal
 - It is queued (*just one*)
 - It will be handled when unblocked

Linux: Signals interface

Service	System call
Send an specific signal	kill
Capture/reprogram a signal handler	sigaction
Block/unblock signals	sigprocmask
Wait UNTIL any event is received (BLOCK)	sigsuspend
Program the automatic sending of signal SIGALRM	alarm

- File `/usr/include/bits/signum.h`
 - Linux implements POSIX interface

Interface: send/capture signals

- Send

```
int kill(int pid, int signum)
```

- signum → SIGUSR1, SIGUSR2, ...

- Capture/Reprogram

```
int sigaction(int signum, struct sigaction *handler,  
struct sigaction *old_handler)
```

- signum → SIGUSR1, SIGUSR2, ...

- handler → *struct sigaction* describes what to do

- old_handler → describes the old behavior

- May be NULL if not interested

A sends signal to B

- Task A sends a signal to B and B runs a specific action when it is received
 - Time of signal reception is unknown

Task A

```
....  
kill( pid, event );  
....
```

Task B

```
int main()  
{  
    struct sigaction hand, old_hand;  
    /* code to initialize hand*/  
    sigaction( event, &hand, &old_hand );  
    ....  
}
```

struct sigaction

```
struct sigaction {  
    void      (*sa_handler) (int);  
    void      (*sa_sigaction) (int, siginfo_t *, void *);  
    sigset_t   sa_mask;  
    int        sa_flags;  
    void      (*sa_restorer) (void);  
};
```

- **sa_handler**: Defines signal management
 - SIG_IGN: Ignore signal at reception
 - SIG_DFL: Default behavior
 - the user function to call when receiving signal
 - The signal identifier is received as the parameter
- **sa_mask**: signals to block during the handler execution
 - The blocked signals are added to the thread mask
 - At least, the captured signal is added
 - At exit from the handler the mask before enter is restored

struct sigaction

```
struct sigaction {  
    void      (*sa_handler) (int);  
    void      (*sa_sigaction) (int, siginfo_t *, void *);  
    sigset_t   sa_mask;  
    int        sa_flags;  
    void      (*sa_restorer) (void);  
};
```

- **sa_flags**: configure behavior
 - 0: Use the default configuration
 - SA_RESETHAND: after executing handler, restore default signal behavior
 - SA_RESTART: if process was blocked in a system_call, restart the blocking system_call

Signals example

```
void main()
{
    char buffer[128];
    struct sigaction hand;
    sigset_t          mask;

    sigemptyset(&mask);
    hand.sa_mask      = mask;
    hand.sa_flags      = 0;
    hand.sa_handler    = int_handler;

    sigaction(SIGINT, &hand, NULL); // Execute int_handler when
                                    // receiving SIGINT

    while(1) {
        sprintf(buffer, "I am working hard\n");
        write(1, buffer, strlen(buffer));
    }
}

void int_handler(int s)
{
    char buffer[128];
    sprintf(buffer, "SIGINT RECIBIDO!\n");
    exit(0);
}
```

Signal mask management

- `sigemptyset`: initialize a mask without blocked signals

```
int sigemptyset(sigset_t *mask)
```

- `sigfillset`: initialize a mask with all signals blocked

```
int sigfillset(sigset_t *mask)
```

- `sigaddset`: add a signal to block to the mask

```
int sigaddset(sigset_t *mask, int signum)
```

- `sigdelset`: remove from the mask a signal

```
int sigdelset(sigset_t *mask, int signum)
```

- `sigismember`: check if a signal is in the mask

```
int sigismember(sigset_t *mask, int signum)
```

Block/unblock signals

- sigprocmask

```
int sigprocmask(int operation, sigset_t *mask, sigset_t *old_mask)
```

- **operation** value:
 - SIG_BLOCK: **Add** signals in *mask* to blocked signals mask of task.
 - SIG_UNBLOCK: **Remove** signals in *mask* from blocked signals mask of task
 - SIG_SETMASK: **Set** *mask* as blocked signals mask

Wait for signals

- sigsuspend

```
int sigsuspend(sigset_t *mask)
```

- Suspend task until a **non ignored signal** is delivered
 - Except: SIGKILL or SIGSTOP
- Change current task's mask by *mask*
 - select signal that will unblock
- After handler execution,
 - signal mask → state before sigsuspend
 - pending signals will be handled

Synchronization: 1st try

- Task B wants to wait for an event from A

Task A

```
.....  
kill( pid, event);  
....
```

Task B

```
int main()  
{  
  sigaction(event, &hand, NULL);  
  ....  
  sigemptyset(&mask);  
  sigsuspend(&mask);  
  ....  
}
```

- ¿What happens if A sends the event BEFORE B arrives to the sigsuspend?
- ¿What happens if B receives another event while in sigsuspend?

Synchronization: 2nd try

Task A

```
.....  
kill( pid, event);  
....
```

- Event is blocked until sigsuspend
- Event is the only accepted to unblock

Task B

```
int main()  
{  
    sigemptyset(&mask);  
    sigaddset(&mask, event);  
    sigprocmask(SIG_SETMASK,&mask,NULL);  
    sigaction(event, &hand, NULL);  
    ....  
    sigfillset(&mask);  
    sigdelset(&mask, event)  
    sigsuspend(&mask);  
    ....  
}
```

Time control

- Program the automatic sending of a SIGALRM
 - The kernel will send it to current process
- `int alarm(int secs);`

```
ret = rem_time;
if (secs == 0) {
    send_SIGALRM = OFF
} else {
    send_SIGALRM = ON
    rem_time = secs,
}
return ret;
```

Time control example

- Task programs a timer for 2 seconds and blocks itself until that time passes

```
int main()
{
    sigemptyset(&mask);
    sigaddset(&mask, SIGALRM);
    sigprocmask(SIG_SETMASK, &mask, NULL);
    sigaction(SIGALRM, &hand, NULL);
    ....
    sigfillset(&mask);
    sigdelset(&mask, SIGALRM);
    alarm(2);
    sigsuspend(&mask);
    ....
}
```

Fork/exec relation

- **FORK: New process**
 - Table of signal handlers is inherited
 - Events are sent to specific processes (PID's), child is a new process → the list of pending signals is reset
 - (pending timers are also reset)
 - The blocked signals mask is inherited
- **EXECLP: Same process, different code**
 - Table of signal handlers is reset
 - Events are sent to specific processes (PID's), process is the same → the list of pending signals is kept
 - The blocked signals mask is kept

Ejemplo 1: gestión de 2 signals (1)



```
void main()
{
    sigemptyset(&mask1);
    sigaddset(&mask1, SIGALRM);
    sigprocmask(SIG_BLOCK, &mask1, NULL);

    trat.sa_flags=0;
    trat.sa_handler = f_alarma;
    sigemptyset(&mask2);
    trat.sa_mask=mask2;
    sigaction(SIGALRM, &trat, NULL);

    sigfillset(&mask3);
    sigdelset(&mask3, SIGALRM);

    for(i = 0; i < 10; i++) {
        alarm(2);
        sigsuspend(&mask3);
        crea_ps();
    }
}
```

```
void f_alarma()
{
}

void crea_ps()
{
    pid = fork();
    if (pid == 0)
        execlp("ps", "ps",
               (char *)NULL);
}
```

Podéis encontrar el código completo en: cada_segundo.c

Ejemplo 2: espera activa vs bloqueo (1)



```
void main()
{
    configurar_esperar_alarma()
    trat.sa_flags = 0;
    trat.sa_handler=f_alarma;
    sigsetempty(&mask);
    trat.sa_mask=mask;
    sigaction(SIGALRM,&trat,NULL);
    trat.sa_handler=fin_hijo;
    sigaction(SIGCHLD,&trat,NULL);
    for (i = 0; i < 10; i++) {
        alarm(2);
        esperar_alarma(); // ¿Qué opciones tenemos?
        crea_ps();
    }
}

void f_alarma()
{
    alarma = 1;
}

void fin_hijo()
{
    while(waitpid(-1,NULL,WNOHANG) > 0);
}
```

```
void crea_ps()
{
    pid = fork();
    if (pid == 0)
        execlp("ps", "ps",
              (char *)NULL);
}
```


Ejemplo 2: espera activa vs bloqueo (2)

Opción 1: espera activa

```
void configurar_esperar_alarma() {  
    alarma = 0;  
}  
void esperar_alarma() {  
    while (alarma!=1);  
    alarma=0;  
}
```



Opción 2: bloqueo

```
void configurar_esperar_alarma() {  
    sigemptyset(&mask);  
    sigaddset(&mask, SIGALRM);  
    sigprocmask(SIG_BLOCK, &mask, NULL);  
}  
  
void esperar_alarma() {  
    sigfillset(&mask);  
    sigdelset(&mask, SIGALRM);  
    sigsuspend(&mask);  
}
```




Distributed memory index

- Signals
- File Descriptors
- Pipes
- Message queues
- Sockets
- Mailboxes


File Descriptors

- In order to send and receive data, a process uses channels to access other devices
- By default, 3 channels
 - 0 → standard INPUT
 - 1 → standard OUTPUT
 - 2 → standard ERROR
- The 3 first entries of the kernel's file descriptor table.
 - A per process table
- When opening a file (or a device) a new entry is added (4, 5, ...)
 - uses the first empty entry in the table



Read and Write

- With these channels you may read/write data
 - (or any other system call using a file descriptor)
- `int read(int fd, char *buffer, int size)`
- `int write(int fd, char *buffer, int size)`



- `write(1, "hello world", 11);`
- `read(0, &buffer, 20);`

IPC


- Connecting the channels of two processes enables a communication between them
- Let's see a couple of ways of doing that...

Distributed memory index

- Signals
- File Descriptors
- Pipes
- Message queues
- Sockets
- Mailboxes


Pipe

- Unidirectional interprocess communication channel
- A read-end and a write-end
 - 2 new file descriptors
- Processes **MUST** be related
 - For example: father/child
- Limited capacity (buffered by kernel)




Pipe interface

- Create a pipe
 - `int pipe(int fd[2])`




Pipe interface

- Duplicate process
 - `fork()`;



Pipe interface


- Select direction
 - `close()`



Pipe behavior

- Blocking device
- Read:
 - **Block** process until data in pipe (not all data requested)
 - If pipe is empty and there are no writers, then EOF (returns 0)
 - If pipe had blocked processes → they are unblocked
- Write:
 - Process write until the pipe is full, then **block** until it is emptied
 - If there is no reader process, writing process will receive SIGPIPE
 - Any other blocked process → unblock
- Unused channels must be closed! Or block.
- This blocking behavior may be used for synchronization

Pipes example



```
void main()
{
    int pipefd[2];
    pipe(pipefd);
    int p = fork();
    if (p == 0) { //Child (Task B)
        close(pipefd[1]);
        int len = read(pipefd[0], buffer, sizeof(buffer));
        write(1, buffer, len);
    }
    else { // Father (Task A == main process)
        close(pipefd[0]);
        write(pipefd[1], "hola", 4);
    }
}
```

Named Pipes

- What happens if you need to communicate 2 processes that are not related?
 - Unnamed pipe can not be used
- A **named pipe** == pipe but:
 - with a name in file system
 - visible globally in the system
 - needs to be opened at both ends before operation
 - opening for read blocks until a writer appear
 - and vice versa

Named pipe creation

- Create a named pipe
 - `#include <sys/stat>`
 - `int mknod(char *path, mode_t mode, dev_t dev)`
or
 - `int mkfifo(const char *pathname, mode_t mode);`
- `mknod("mipipe", S_IFIFO|S_IRUSR|S_IWUSR, 0)`
or
- `mkfifo("mipipe", S_IRUSR|S_IWUSR)`

mipipe



Access named pipe

- Open named pipe
 - `int open(const char *pathname, int flags);`
 - Process A
 - `fd = open("mipipe", O_WRONLY)`
 - `write(fd, "hola", 4);`
 - Process B
 - `fd = open("mipipe", O_RDONLY)`
 - `read(fd, buffer, sizeof(buffer))`



Avoid blocking

- The blocking on the pipe can be avoided
 - `O_NONBLOCK` flag
- At *open* or at creation through *fcntl*
 - `#include <fcntl.h>`
 - `int fcntl(int fildes, int cmd, ...);`
 - `open("mipipe", O_RDONLY | O_NONBLOCK)`
or
 - `fcntl(fd, F_SETFL, O_NONBLOCK)`

Non blocking behavior


- Open named pipe:
 - Read only: returns the file descriptor wo block
 - Write only: returns error (`errno==ENXIO`)
- Read on an empty pipe with writers:
 - returns error (`errno == EAGAIN`)
- Write on a full pipe:
 - returns error (`errno == EAGAIN`)

Distributed memory index

- Signals
- File Descriptors
- Pipes
- Message queues
- Sockets
- Mailboxes

Message queues

- Solve some of the limitations of pipes
 - Granularity is message not bytes
 - Messages are atomic



Message queues interface

```
#include <fcntl.h>          /* For O_* constants */
#include <sys/stat.h>        /* For mode constants */
#include <mqueue.h>
```

- `mqd_t mq_open(const char *name, int oflag);`
- `mqd_t mq_open(const char *name, int oflag, mode_t mode, struct mq_attr *attr);`
- `int mq_send(mqd_t mqdes, const char *msg_ptr, size_t msg_len, unsigned int msg_prio);`
- `ssize_t mq_receive(mqd_t mqdes, char *msg_ptr, size_t msg_len, unsigned int*msg_prio);`
- `int mq_close(mqd_t mqdes);`
- Link with `-lrt`.

POSIX Message queues

- Create and open using `mq_open(3)`
→ returns a message queue descriptor (`mqd_t`)
- Transfer messages to and from a queue using `mq_send(3)` and `mq_receive(3)`
- When a process has finished using the queue, it closes it using `mq_close(3)`
- When the queue is no longer required, it can be deleted using `mq_unlink(3)`

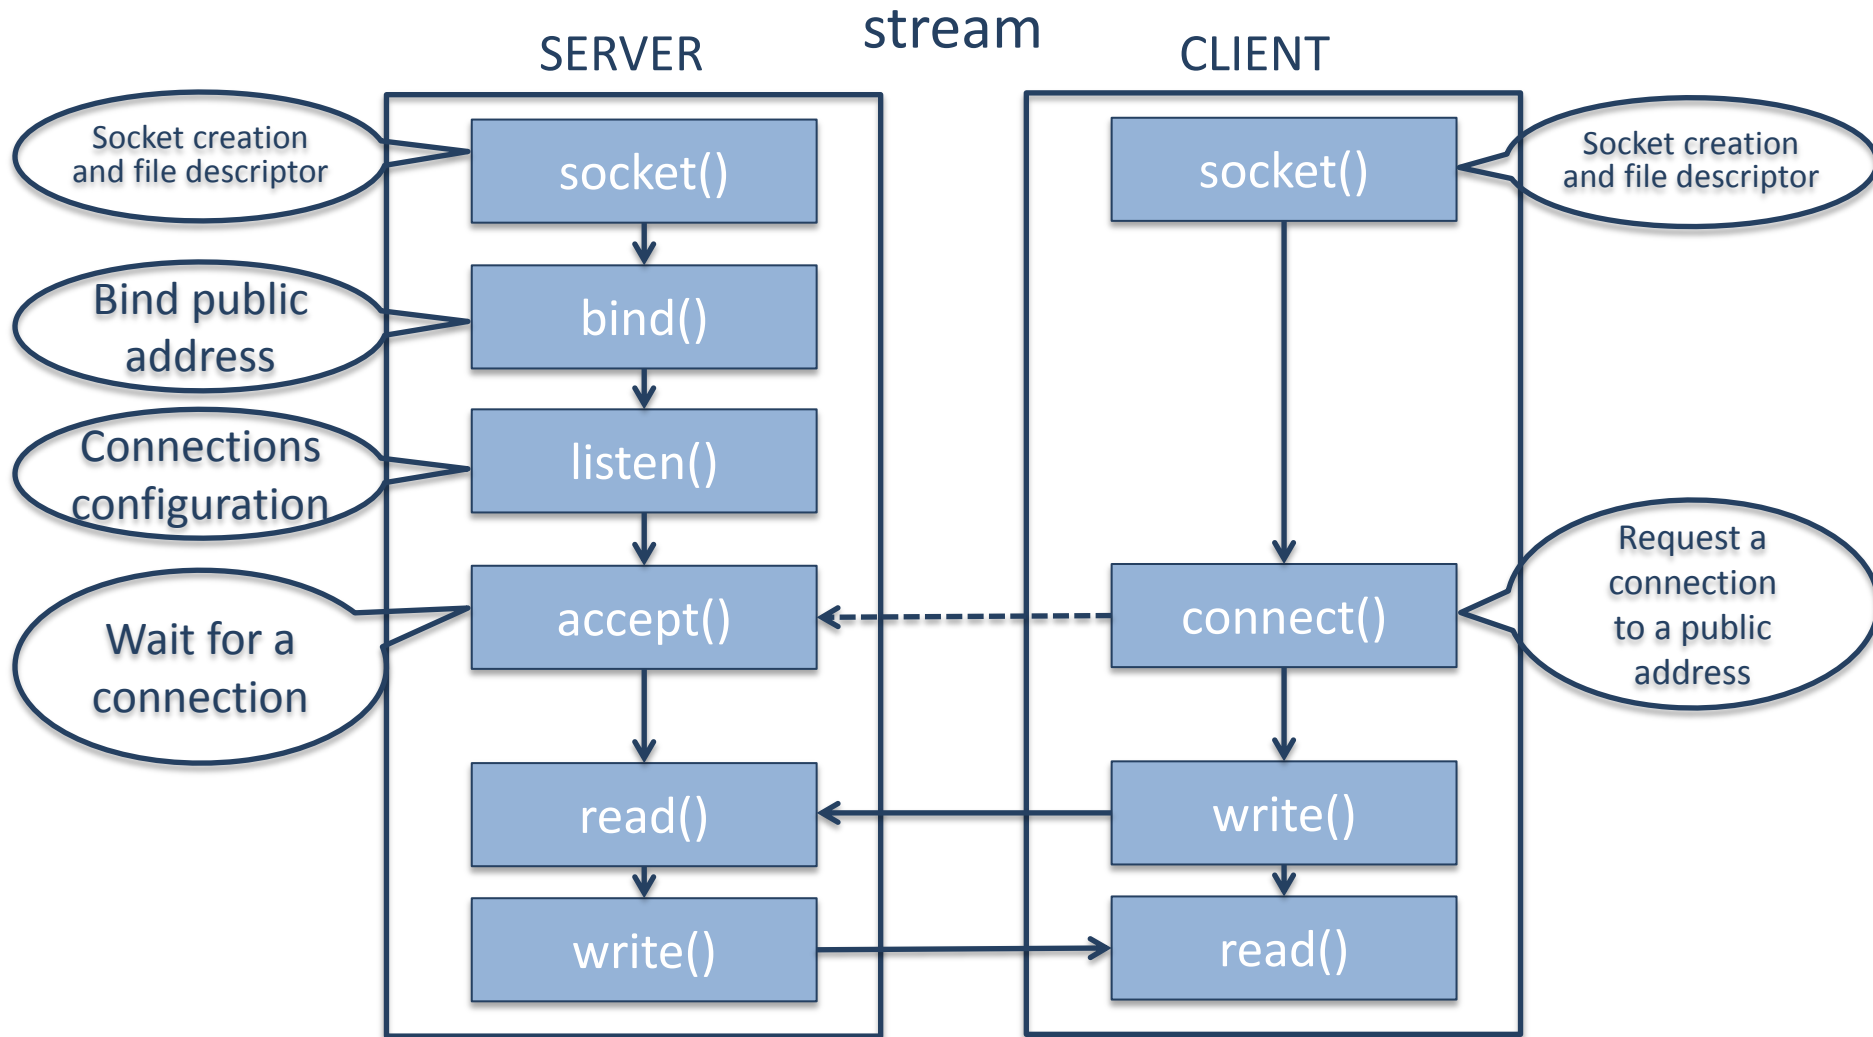
Distributed memory index

- Signals
- File Descriptors
- Pipes
- Message queues
- Sockets
- Mailboxes

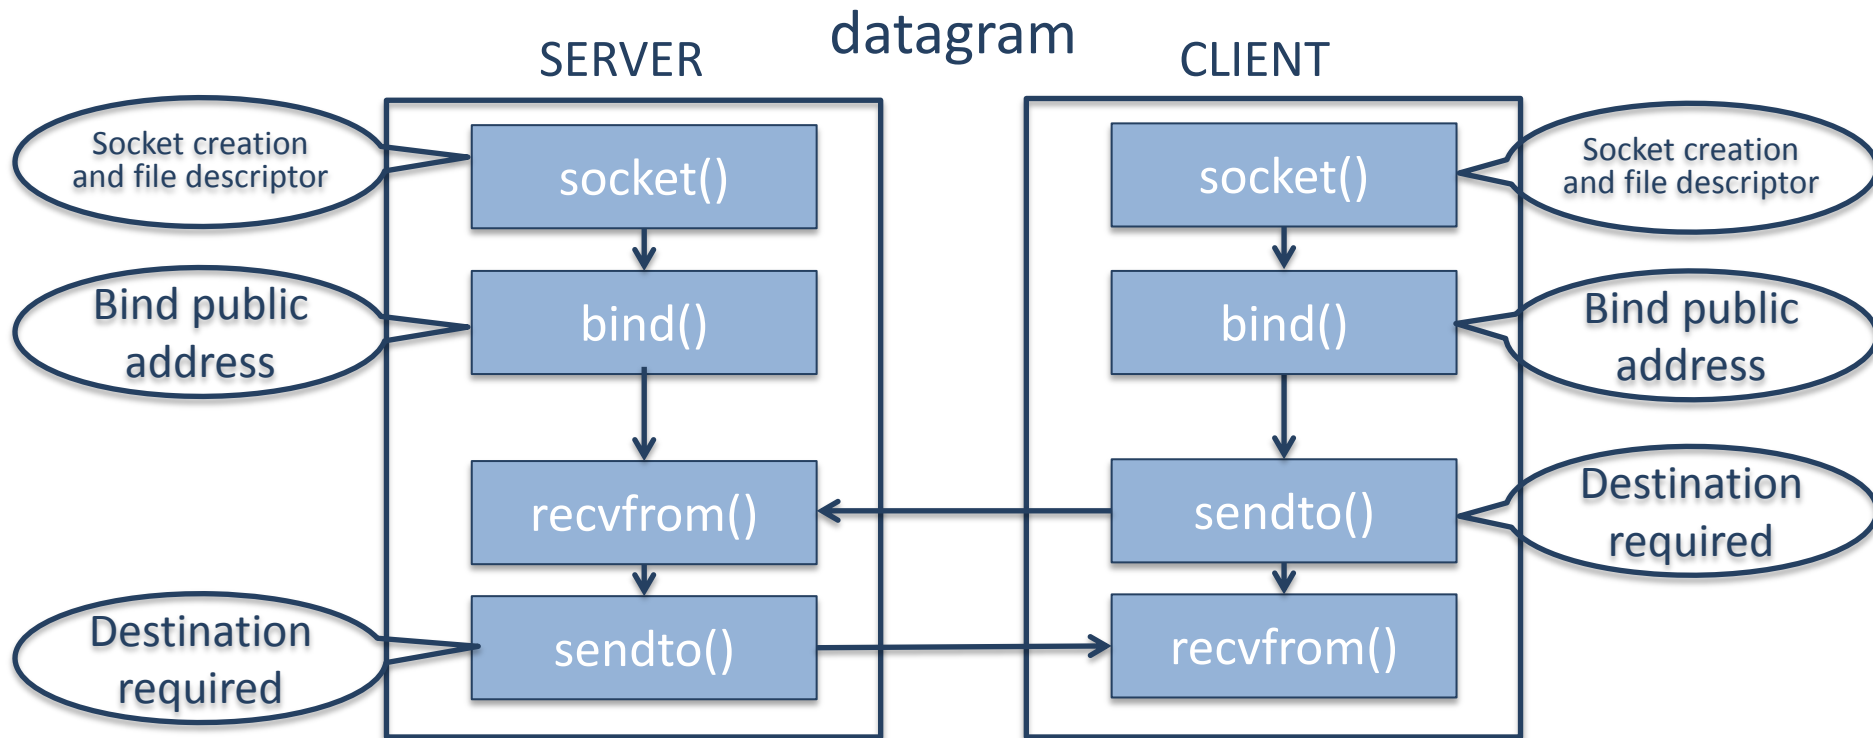
Sockets

- Bidirectional channel between processes
 - Local
 - Remote (through network)
- The following must be defined:
 - Namespace (Unix, INET, ...)
 - Communication type (stream or datagram)
 - Communication protocol (TCP, UDP, ...)

Sockets: Stream communication



Sockets: Datagram connection



Endianness

- When transmitting data from one machine to another, they may have different ways to store complex types
 - Big Endian
 - Little Endian
- In order to transfer these types you need to translate them
 - 1) from host to network (htons, htonl)
 - 2) from network to host (ntohs, ntohl)

Create socket

- `#include <sys/types.h>`
- `#include <sys/socket.h>`
- `int socket (int af, int type, int proto)`
 - `af`: address family(namespace to use)
 - `AF_UNIX`
 - `AF_INET`
 - `type`: connection type
 - `SOCK_STREAM`
 - `SOCK_DGRAM`
 - `proto`: Use 0 to automatically use the most appropriate
 - Returns the *file descriptor* associated to the socket or -1 if error

Bind address

- `#include <sys/socket.h>`
- `int bind (int fd,`
 `struct sockaddr *addr, int addr_size)`
 - `fd`: the one returned by `socket`
 - `addr`: address to connect to this socket
 - `addr_size`: socket address size in bytes
 - returns 0 if ok or -1 if error
- *addr* uses a generic type (`struct sockaddr`) but depending on the socket's namespace you must use the specific type
 - `AF_UNIX`: nombre de fichero [`struct sockaddr_un`]
 - `AF_INET`: IP+puerto [`struct sockaddr_in`]

Socket addresses

- Generic structure for addresses

```
#include <sys/un.h>
struct sockaddr {
    unsigned short sa_family;    // address family, AF_XXX
    char           sa_data[14]; // 14 bytes of protocol address
};
```

- This is just a “container” for the specific addresses... DO NOT USE directly!

Socket addresses

- Socket namespace: PF_UNIX: path to a file

```
#include <sys/un.h>
struct sockaddr_un {
    sa_family_t sun_family;    /* 1 byte == AF_UNIX */
    char sun_name[UNIX_PATH_MAX]
}
```

- Socket namespace: PF_INET: ip address + port

```
#include <sys/netinet/in.h>
struct sockaddr_in {
    sa_family_t    sin_family; /* 1 byte == AF_INET */
    struct in_addr sin_addr;    /* 4 bytes */
    in_port_t      sin_port;    /* 1 byte */
    char           sin_zero [8]; /* not used, must be 0 */
}
```

- sin_addr: INADDR_ANY represents machine's IP address running the code
- sin_port : if 0, system assigns an empty one; otherwise ensure that it is not used.
 - Use ports > 5000
 - Use **Network codification**

Getting socket addresses

- IP → binary

```
struct sockaddr_in sa; // IPv4
inet_pton(AF_INET, "10.12.110.57", &(sa.sin_addr));
```

- binary → IP

```
char ip4[INET_ADDRSTRLEN]; // space to hold
                             // the IPv4 string
struct sockaddr_in sa; // pretend this is
                        loaded with a valid address..

inet_ntop(AF_INET, &(sa.sin_addr),
           ip4, INET_ADDRSTRLEN);
```

Request a connection

- `#include <sys/socket.h>`
- `int connect(int fd,
 struct sockaddr *addr, int addr_size)`
- If the server may not attend request
 → return -1
- Else, **block** until server accepts the request
 - If `O_NONBLOCK` enabled → no block and return -1

Automating address resolution

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
```

```
int getaddrinfo (const char *node,
                  const char *service,
                  const struct addrinfo *hints,
                  struct addrinfo **res);
```

- node: host name or IP
 - service: port number or service name
 - hints: specifies criteria for selecting the socket address structures returned in the list res
 - res: linked list with the results
- Returns 0 if it succeeds

Struct addrinfo

```
struct addrinfo {  
    int             ai_flags;        // AI_PASSIVE, AI_CANONNAME, etc.  
    int             ai_family;       // AF_INET, AF_INET6, AF_UNSPEC  
    int             ai_socktype;     // SOCK_STREAM, SOCK_DGRAM  
    int             ai_protocol;     // use 0 for "any"  
    size_t          ai_addrlen;      // size of ai_addr in bytes  
    struct sockaddr *ai_addr;         // struct sockaddr_in or _in6  
    char            *ai_canonname;   // full canonical hostname  
  
    struct addrinfo *ai_next;        // linked list, next node  
};
```

```
struct addrinfo hints;  
memset(&hints, 0, sizeof hints); // make sure the struct is empty  
hints.ai_family = AF_UNSPEC;      // don't care IPv4 or IPv6  
hints.ai_socktype = SOCK_STREAM; // TCP stream sockets  
hints.ai_flags = AI_PASSIVE;      // fill in my IP for me  
...
```

Wait for a connection request

- `#include <sys/socket.h>`
- `int accept(int fd,`
 `struct sockaddr *addr, int *addr_size)`
 - `fd`: the one returned by `socket`
 - `addr`: will contain the address requesting the connection
 - `addr_size`: address size in bytes
 - returns
 - a new file descriptor to use for connection
 - or -1 if error
- If pending connections → accept first one
- Otherwise, **block** until a new request
 - unless the flag `O_NONBLOCK` enabled
 - no block and return error

Sockets example

- Beej's Guide to Network Programming
<http://beej.us/guide/bgnet>

Distributed Memory Summary

- There is a need to communicate and synchronize processes
 - Without shared memory
- Different mechanisms
 - Events: Signals
 - IPC
 - Pipes(Unnamed & Named)
 - Message Queues
 - Sockets

References

- [Harder18] Harder, Douglas W et al. 2018. “A practical introduction to real-time systems for undergraduate engineering”
- [Downey08] Downey, Allen B. 2008. “The Little Book of Semaphores”. Green Tea Press
- [JPL09] Jet Propulsion Laboratory. 2009. “JPL Institutional Coding Standard for the C Programming Language” California Institute of Technology.
- [Burns09] Burns, Alan and Wellings, Andy. 2009. “Real-Time Systems and Programming Languages”
- [Bertolotti12] Bertolotti, Ivan and Manduchi, Gabriele. 2012. “Real-time embedded systems. Open-source Operating systems perspective”
- [Stevens98] Stevens, Richard. 1998. “Unix Network Programming”, volumes 1-2.